

Supervised Machine Learning

Supervised Learning

Your "training" dataset is composed of examples of labeled examples:

- - features
- - labels

A supervised learning model learns to predict the label from the features .



Supervised Learning Examples

Credit Card Fraud Detection

- **Features:** Vendor, location, time, distance from last transaction
- **Labels:** Chargebacks on previous transactions

Question: What kind of supervised learning?

1. Regression
2. **Binary Classification**
3. Multiclass Classification

[Card Fraud Prevention] Activity On Your Debit or ATM Card On 12/28/2019 [MAIL ID:4435446]

CF

Chase Fraud Alert <admin@vagar.com>

Saturday, December 28, 2019 at 8:00 AM

Graham, Jefferson

Show Details

Email not displaying correctly? [View it](#) in your browser

CHASE

Dear Customer,

We're letting you know that we've detected some unusual activity on your card on 12/28/2019. For your security, please verify the following transaction(s) so that you can continue to use your card

Do you recognize all of these transaction(s)?

Approved transaction at SQC*CASH APP for \$224.49 on 12/28/2019

Declined transaction at TOP UP B.V. for \$624.11 on 12/28/2019

Approved transaction at BESTBUY for \$124.59 on 12/28/2019

YES, I recognize all of these transactions

YES will make your card immediately ready to use again

NO, I don't recognize one or more of these transactions

NO will allow you to complete the verification process and file a fraud claim in Online or Mobile Banking

Please do not reply to this automatically generated message. If you have any questions, please call us at the number located at the top of your statement.


Supervised Learning Examples

Stock Market Prediction

- **Features:** Stock price from Feb 1st to March 1st
- **Labels:** Stock price on March 7th.

Question: What kind of supervised learning?

1. Regression
2. Binary Classification
3. Multiclass Classification



Supervised Learning Examples


Classification

- **Email Spam Filters**
 - Features: Words, sender, links
 - Label: Spam or Not Spam
- **Face ID / Fingerprint Unlock**
 - Features: Facial/fingerprint data
 - Label: You or Not You
- **Letter Recognition**
 - Features: Pixel values of images
 - Label: A-Z, 0-9, etc.


Regression

- **Weather Forecasting**
 - Features: Pressure, humidity, wind
 - Label: Temperature/rainfall amount
- **Uber/Lyft Pricing**
 - Features: Distance, time, demand
 - Label: Trip cost
- **YouTube View Count Predictions**
 - Features: Title, thumbnail, creator stats
 - Label: Expected views








































Real-World Example

Outer Wall Thickness of an extruded vinyl profile is measured and recorded manually using a cut profile and a pair of calipers once every 12 hours.

- If this wall is **too thin**, the profile will create failure points.
- If this wall is **too thick**, the profile will be too heavy and expensive.
- Customer is **losing \$400K/mo** in overage.

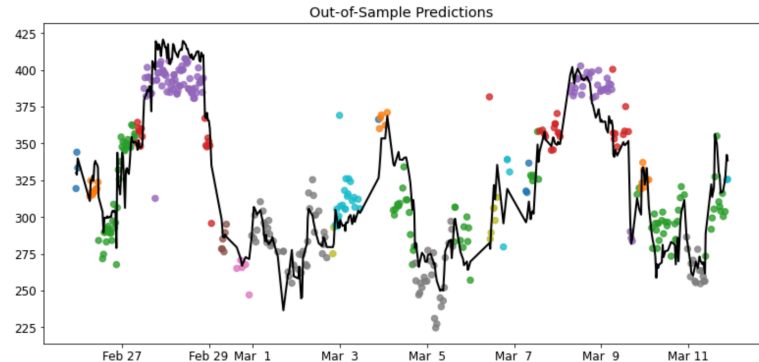
Metrics

- **Main Feed Speed:** Speed at which substrate material feeds into the main extruder.
- **Main Drive Amps:** Electrical load on the main extruder's motor.
- **Puller Speed:** Speed at which the product is pulled from the die.
- **Die and Barrel Zone Temperatures:** Temperature at the die and extruder barrel.

Model

- Lasso Regression scored against a 60-minute rolling window
- Pre-aggregated data into 5-minute buckets

Dataset	Accuracy (>90%)	Correlation (>70%)	P95 error (<3 sigma)
Within-Sample	96%	94%	1.99 sigma
Out-of-Sample	95%	93%	1.73 sigma



Exercise: Predicting Passengers

bigd103.link/linear-regression

Supervised Classification

Exercise

On average, which of the following Pokemon cards does the most damage? Write your answer in chat.



A Brief Interlude into Probability

Expected Value

Expected Value is a way to measure the average outcome of a random variable.

For example, “What’s the expected value of a standard die roll?” you’d add its possible outcomes and divide by 6:

$$\mathbb{E}[\text{die}] = \frac{1 + 2 + 3 + 4 + 5 + 6}{6} = 3.5$$

Using this, we can answer:

$$\begin{aligned}\mathbb{E}[\text{Quick Attack}] &= \left(\frac{1}{2}\right) \times 30 + \left(\frac{1}{2}\right) \times 10 \\ &= 15 + 5 \\ &= 20\end{aligned}$$



Multiplication Rule

The probability of something happening is denoted as:


$$P(\text{event})$$

Two independent events have a joint probability:


$$P(A \cap B) = P(A) \times P(B)$$

For example, if the probability of flipping a coin on heads once is $1/2$ than the probability of flipping it twice and getting heads both times is:

$$P(\text{heads}) \times P(\text{heads}) = 1/2 \times 1/2 = 1/4$$












Logistic Regression







Decision Boundary


The decision boundary is defined where the model predicts $P(y = 1 \mid x) = 0.5$.


Equivalently, $\sigma(z) = 0.5 \rightarrow z = 0$.

Loss Function and Optimization

Uses the **Cross-Entropy Loss** (also called Log Loss).

Parameters (w_i) are optimized using **gradient descent** to minimize this loss.








Training a Supervised Learning Model
















What is an SVM?

Support Vector Machines are powerful classifiers that find the **optimal decision boundary** between classes.


Key idea: Don't just find any line that separates the classes - find the line with the **maximum margin**.

Margin: The distance between the decision boundary and the nearest data points from each class.











The Kernel Trick


Problem: What if the data isn't linearly separable?

Solution: The kernel trick allows SVMs to:

1. Map data to a higher-dimensional space
2. Find a linear boundary in that space
3. Which appears non-linear in the original space!

Common Kernels

- **Linear:** no transformation
- **RBF:** Radial Basis Function
- **Polynomial:** polynomial transformation

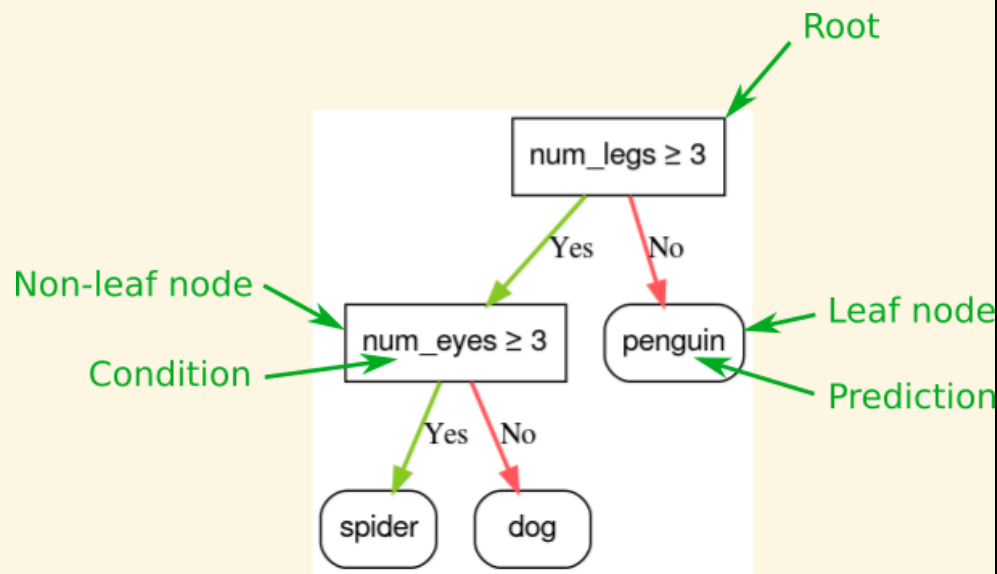




What is a Decision Tree?

A decision tree is a flowchart-like structure used for classification and regression tasks.


It recursively splits the dataset into subsets based on feature values, forming a tree of decisions.


























Comparing Classifiers

Algorithm	Pros	Cons
Logistic Regression	Fast, probabilistic, interpretable	Linear only
SVM	Powerful, kernels, works in high-D	Slow, hard to interpret
Decision Trees	Interpretable, non-linear, fast	Can overfit easily

In practice: Try multiple algorithms and see what works best for your data and problem!

Machine Learning in Production

